

Exploring Extractive and Abstractive Approaches for Multi-Document Summarization: An End-to-End System with Benchmarking and Error Analysis

Anna Batra, Sam Briggs, Junyin Chen, Hilly Steinmetz

Department of Linguistics, University of Washington
{batraa, briggs3, junyinc, hsteinm}@uw.edu

Abstract

As the number of online publications continues to grow at a rapid pace, there is a pressing need for a valuable tool that can automatically generate summaries from multiple news articles. While most existing works in this field have focused on single document summarization, our paper explores multiple approaches to multi-document summarization. These approaches include both extractive and abstractive methods. We build end-to-end systems using the selected methods. Finally, we will benchmark selected methods and perform an error analysis to evaluate their effectiveness.

1 Introduction

Text summarization is a process of generating summaries that are both accurate and concise from one or more input documents. It is an important task in Natural Language Processing and its applications are growing due to the increasing demand for concise and easily-understood content. There are currently two approaches, extractive or abstractive summarization. Extractive summarization pulls key phrases from the source document and combining them to make a summary. Abstractive summarization generates new sentences that capture the main ideas of the source documents. A well-written abstractive summary includes the main information from the source and is expressed in fluent language.

Creating a well-organized summary that comprehensively covers a news event while avoiding repetition is a challenge when summarizing from multiple documents. The input documents may have varying focuses and viewpoints on the event.

In the past, summarization was tackled by using non-neural methods by turning it into a binary integer linear programming problem (ILP) or by employing the LexRank algorithm to rank sentences automatically.

Neural methods for text summarization have recently advanced and have mostly been used for

single-document summarization (SDS) and headline generation. Fabbri et al. (2019) create the Multi-News dataset, the first large-scale Multi-document summarization (MDS) news dataset, as previous MDS datasets such as TAC 2011 (Owczarzak and Dang, 2011) as less than 100 document clusters. However, current popular language model for summarization, such as BART (Lewis et al., 2019), T5 (Raffel et al., 2019), PEGASUS (Zhang et al., 2019), optimized for single document summarization.

In our paper, we have made the following contributions: We evaluated the performance of LexRank and topic clustering algorithms by using TF-IDF, Word2Vec, and DistilBERT to create document vectors. We also evaluated the performance of PEGASUS when dealing with MDS dataset. Finally, we build end-to-end systems to incorporate various methods and benchmark our system’s performance.

2 Related Works

Traditionally, non-neural methods have been applied to multi-document summarization task, which can be either extractive (Carbonell and Goldstein, 1998; Radev et al., 2000; Erkan and Radev, 2004; Mihalcea and Tarau, 2004; Haghighi and Vanderwende, 2009) or abstractive methods (McKeown and Radev, 1995; Barzilay et al., 1999; Ganesan et al., 2010). However, in recent years, neural methods have emerged as a promising alternative for text summarization, with both extractive (Nallapati et al., 2016b; Cheng and Lapata, 2016; Narayan et al., 2018) or abstractive methods (Chopra et al., 2016; Nallapati et al., 2016a; See et al., 2017; Paulus et al., 2017; Cohan et al., 2018; Celikyilmaz et al., 2018; Gehrmann et al., 2018). Now, a significant boost in performance for both natural language understanding and text generation tasks has been achieved by fine-tuning Transformer-based sequence models that have been pre-trained with

much larger external text corpora.

3 System Overview

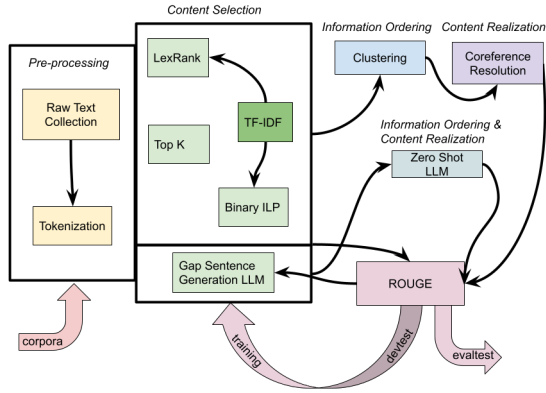


Figure 1: The base system architecture from preprocessing the data to content realization

After pre-processing steps including raw text collection and tokenization of paragraphs, sentences, and words we build off of four content selection models. One is our baseline, TopK (an extractive approach). Two use an extractive approach and build off of a shared TF-IDF class: Integer Linear Programming & LexRank. The fourth content selection model uses an abstractive approach, using gap sentence generation on a large language model. It has a direct connection with the ROUGE score as it uses it to figure out the importance of sentences during training.

TopK, ILP, and LexRank all use clustering and coreference resolution as information ordering and content realization methods respectively. The LLM uses zero shot learning to improve upon information ordering and content realization together.

4 Approach

4.1 Data Pre-processing

4.1.1 Accessing Input Data

We used the data from the TAC 2009 (Owczarzak and Dang, 2009), 2010 (Owczarzak and Dang, 2010), 2011 (Owczarzak and Dang, 2011) Shared Task data. To process the given XML files to retrieve the set of document names in docSetA, we used `xml.etree.ElementTree`. We then figured out which corpus data path matched each document in docSetA.

In order to read in these AQUAINT and AQUAINT2 files that are organized differently, we

used `lxml.etree` as this can parse non-XML compliant files. We found there are three different organization methods that were used, and we made sure to process each kind uniquely. From these files, we parsed the the headline, time, and raw text paragraphs.

4.1.2 Tokenization

After getting the raw text, we used two different tokenization methods: spaCy and NLTK.

From spaCy 2.0 we used the English model tokenizer “en_core_web_sm”. We decided on spaCy 2.0 since it fits the python version (Python 3.6) on Patas (the virtual machine we used). We decided not to tokenize the raw text into sentences first. Instead, we just ran a word tokenizer on the raw text. The spaCy tokenizer utilizes a transformer under the hood. We then realized that there was a sentence tokenizer, but since this tokenization method uses transformers on the CPU, it actually takes a really long time to run. Therefore, we decided to leave it as it is and turn to a rule-based tokenizer explained in the next method.

For NLTK, we used the English model tokenizer “tokenizers/punkt/english.pickle”. We first ran a sentence tokenizer, and then for each tokenized sentence, we then ran a word tokenizer. This gave us a list of sentences, where each sentence contains tokenized words. The NLTK tokenizer is significantly faster than spaCy’s tokenizer and gives us the correct output we need per D2.

4.1.3 Sentence Embeddings

We wanted to investigate whether semantic information can improve the performance of various algorithms used generate summaries. We compared the performance of the LexRank and topic clustering algorithms by utilizing TF-IDF, Word2Vec (Mikolov et al., 2013), and DistilBERT (Sanh et al., 2019) to generate document vectors. For the pre-trained word2vec model, we used the continuous bag of words (CBOW) model trained on the Google News corpus.

Creating vectors from each document involve different pooling methods. The TF-IDF vector was constructed by using either calculating TF-IDF for each document, with IDF values obtain either from the document set or data set. Word2vec vectors were created by averaging each word vector (obtained from the pretrained model). DistilBERT vectors were created by obtaining the final hidden states of the pretrained model, and averaging these

final states.

4.2 Content Selection

4.2.1 Baseline: Top K

For our baseline we implemented taking the first k sentences of the first document in the docset. If the first sentence is too long (over 100 words), we keep skipping the first few sentences until we find one sentence less than 100 words. Then, we continue adding more sentences one by one until the next sentence makes the summary over 100 words. This means possibly we could end up with a one sentence summary if the proceeding sentence already makes it over 100 words.

4.2.2 TF-IDF

To obtain the importance of an n -gram in a given document set, we used the “term-frequency, inverse document frequency ($tf \cdot idf$)” metric. To calculate, we used a few different formulas with a few different parameters. The parameters were as follows:

- Document_level: Whether to treat each sentence as a document, or the entire docset as a document.
- N-gram: Whether to treat each term as a unigram, bigram, or trigram. Padding was incorporated here for both start of sentence and end of sentence tokens, using `nlk.util.ngrams`.
- Eliminate_punctuation: Whether to include punctuation or not.
- Casing: Whether to lowercase all letters, or maintain original capitalization
- log: Whether to use logged equations or not (see equations below)
 - log_base: If logged equations are used, what base to use
- smoothing: Whether to smooth tf and idf , or not
 - tf_delta: Which δ_1 to use or tf ftf smoothing
 - idf_delta: What δ_2 to use for idf if smoothing

One difference than normal tf - idf is that we used tf at a different level of document than idf . For LexRank we used a sentence level document, while allowing the idf to span over the entire dataset.

For ILP, we used a docset level document, while allowing the idf to span over the entire dataset. This will help to make frequent words insignificant and help located the more important words for the sentence/docset.

Using the logarithmically scaled, add δ_1 smoothed tf , and we used an add δ_2 smoothed idf to weight each term in the document set (Seki, 2003).

Given all the training data D , an n -gram t , and a document set $d \subseteq D$, we calculated the logged term-frequency, inverse document frequency ($tf \cdot idf$) as follows:

First, we let:

$$f_{t,d} = \text{count}(t) \text{ for } t \in d \quad (1)$$

$$n_t = |\{d \mid t \in d, d \in D\}| \quad (2)$$

If logged, we calculate as follows:

$$tf \cdot idf(t, d, D) = tf(t, d) \cdot idf(t, d, D) \quad (3)$$

$$tf(t, d) = \log(\delta_1 + f_{t,d}) \quad (4)$$

$$idf(t, D) = \delta_2 + \log\left(\frac{N}{\delta_2 + n_t}\right) \quad (5)$$

If not logged, we calculate as follows:

$$tf \cdot idf(t, d, D) = tf(t, d) \cdot idf(t, d, D) \quad (6)$$

$$tf(t, d) = \delta_1 + f_{t,d} \quad (7)$$

$$idf(t, D) = \delta_2 + \frac{N}{\delta_2 + n_t} \quad (8)$$

If not smoothed, δ_1 and δ_2 effectively become 0.

4.2.3 Binary Linear Programming

For the ILP task, previous work has been done by Gillick et al. (2008) and Luo et al. (2018). In line with both Gillick et al. (2008) and Luo et al. (2018), we used n -grams for “concepts”. We tried unigrams, bigrams, and trigrams (exclusively) as concepts. Unlike Luo et al. (2018) who used *term-frequency* for their concept weights, and Gillick et al. (2008) who used *document frequency* for their concept weights, we looked to combine the two weighting methods and used the *term-frequency inverse document frequency* (tf - idf) of n -grams as calculated in section 4.2.2. For the formulation of the ILP, we used the objective function, constraints, and binary variables as proposed in Gillick et al. (2008).

For notation, we take a bag of sentences and bag of concepts approach. We call the given set of sentences Y which constitute the given a document set, and the set of concepts Z which constitute the given document set.

We use $y_j \in Y$ for sentence j and we use $z_i \in Z$ for concept i . We also let y_j and z_i be indicator functions, indicating whether to include or exclude sentence y_j and concept z_i respectively from the summary, and thus y_j and z_i can only take on values of 0 or 1.

We use $A_{i,j}$ to denote the indicator function $\mathbb{1}_{z_i \subseteq y_j}$, i.e. $A_{i,j} = 1$ if concept z_i appears in sentence y_j , 0 otherwise. We use the weight $w_i \in \mathbb{R}$ where weight w_i is the corresponding weight for "concept" z_i . We also have a maximum term summary length L . If we have N sentences in the optimal summary, and M sentences total in the document set, we can then formulate the optimization problem as follows:

$$\text{maximize}_{y,z} \sum_{i \in Z} w_i z_i \quad (9)$$

$$\text{Subject to } \sum_j^M A_{i,j} y_j \geq z_i, \forall i \in Z \quad (10)$$

$$A_{i,j} y_j \leq z_i, \forall i, j \in Z \times Y \quad (11)$$

$$\sum_j^N l_j y_j \geq L \quad (12)$$

$$y_j \in \{0, 1\} \quad (13)$$

$$z_i \in \{0, 1\} \quad (14)$$

We see that (Eq. 9) is the objective function where we are trying to maximize the total weight of the concepts chosen for the summary in an attempt to extract the most important concepts in a document set. (Eq. 10) ensures that a concept is included in the optimal summary if and only if there is a sentence that is selected for the optimal summary that contains said concept. (Eq. 11) ensures that all concepts in sentence y_j are included if y_j is included in the optimal summary. (Eq. 12) ensures that the given summary remains under the maximum sentence length L . For each "concept" z_i , we tested unigrams, bigrams, and trigrams. For the corresponding weight w_i for each concept, we used the tf-idf score of the unigram, bigram, or trigram z_i (exclusive) as calculated in section 4.2.2.

Although there is no explicit redundancy

checker, we see that implicitly, redundancy is kept to a minimum because of the formulation of the ILP problem—each concept will appear in the optimal summary only once.

We found that the best combination of hyperparameters to pass in to our calculation of tf-idf was to use calculate tf-idf for unigrams, using a logged tf-idf with δ_1 close to 0 (in our case $\delta_1 = 0.01$), a δ_2 close to 1 (our best case uses $\delta_2 = 0.7$). We also found that eliminating punctuation and lowercasing all tokens yielded the best results. We also found that removing sentences with less than around 25 tokens (after being tokenized by `nlk.word_tokenize`) yielded the best ROUGE scores.

4.2.4 LexRank

We also implemented the LexRank algorithm described in [Erkan and Radev \(2004\)](#). LexRank is an adaption of the PageRank algorithm ([Page et al., 1999](#)), and was proposed as an alternative to centroid-based approaches. LexRank leverages relationships between documents by creating a weighted graph that connects sentences. Relating the sentences to one another has the advantage of (1) dampening the effect of high IDF scores of rare words (when using TF-IDF vectors) and (2) formalizing a preference for more informative (or more connected) sentences.

The LexRank algorithm treats each sentence as a document. It compares sentence vectors to construct a weighted graph of the relationships between sentences in a document set. [Erkan and Radev \(2004\)](#) obtains sentence vectors using TF-IDF (without smoothing); however, sentence vectors can be obtained using a number of methods (see 4.1.3).

Sentences are compared to one another are related to one another using the cosine similarity measure:

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\|_2 \times \|\mathbf{y}\|_2} \quad (15)$$

A similarity matrix can then be constructed by calculating similarity scores across all sentences in the document. Unlike [Erkan and Radev \(2004\)](#), we do not calculate the similarity between instances of the same sentence, since we found these high scores to hurt its performance. We believe the decline in performance might pertain to lowering other similarity scores after the matrix is normalized by rows.

Using this similarity measure, we created a similarity matrix between sentences in the document,

which also functioned as a weighted graph. Per Erkan and Radev (2004), values with low similarities scores are discarded and self-connections between nodes. The matrix satisfies the properties of a stochastic matrix, allowing us to use the power method to estimate the eigenvalue of the matrix. We initialize the centrality vector as $\mathbf{p} = \frac{1}{N}\mathbf{p}$, where N is the number of documents. We then apply the following update to \mathbf{p} :

$$\mathbf{p} = [d\mathbf{U} + (1 - d)\mathbf{B}]^T \mathbf{p} \quad (16)$$

where \mathbf{U} is a square matrix of size $[N \times N]$ with values equal to $1/N$ and \mathbf{B} is the adjacency matrix of the graph.

Experiments were conducted to investigate whether LexRank performed best with TF-IDF, Word2vec, or DistilBERT vectors. A comparison table can be found in Appendix B.

4.2.5 Gap sentence generation

We used the gap sentence generation method introduced in Zhang et al. (2019). Based on the finding when Zhang et al. were training the Pegasus based model, we will select the top m sentences as gap sentences without replacement from a document based on importance score. The importance score is calculated based on the ROUGE score one sentence get comparing to the remaining sentences in one document as in Algorithm 1.

Algorithm 1 Independent sentence selection

```

1:  $D := \{x_i\}_n \leftarrow$  sentences in alldocument
2:  $S := \emptyset$ 
3:  $I \leftarrow$  list contains index from 0 to  $n$ 
4: for  $j \leftarrow 1$  to  $n$  do
5:    $s_i := \text{rouge}(x_i, D \setminus \{x_i\})$ 
6:    $S := S \cup \{s_i\}$ 
7:  $I := \text{sort}(I)$  Based on the value in  $S$ 
```

To improve the system, when we are processing the training data, we calculate the ROUGE score based on the average ROUGE score when comparing the selected sentence with each of the gold summary. We only calculate the ROUGE score based on the rest of the documents when we are processing the test and validation data.

We only mask the top thirty percent of the sentences as Zhang et al. finds out achieves relatively high performance without sacrifice training efficiency.

4.3 Information Ordering

4.3.1 Topic Clustering

This algorithm tries to order topics in the order that they are most likely to appear in a document. The intuition being that similar sentences form topics, and topics must be ordered in the original documents in a cohesive manner. This algorithm tries to recreate this cohesive topic order. This algorithm is heavily modified from the 'Augmented Ordering Algorithm' presented in "Barzilay et al. (2002). Barzilay groups similar themes from different documents into blocks, and then orders the blocks by timestamp. However, in our algorithm, we instead put all of the documents together and consider it one document, clustering similar sentences/topics of information into themes. There are no senses of block here as discussed in "Barzilay et al. (2002).

Similarity between sentences is determined by the similarity of sentence embeddings. To group topics, we group similar sentences over a whole document set. We calculate sentence as calculated in section 4.1.3 using TFIDF embeddings for each word in a sentence. We chosen between TFIDF, Word2Vec, and DistilBERT as a hyperparameter for sentence embeddings. Please see Table 1 for an error analysis between them.

We then use sentence embeddings to create topic clusters to group similar sentences. To group sentences, we used K-means clustering. To run K-means clustering, we used sklearn.cluster.KMeans with 8 clusters, and the following parameters:

```

kmeans = KMeans(
    n_clusters=8, init='k-means++',
    n_init=10, max_iter=300,
    tol=0.0001, verbose=0,
    random_state=None, copy_x=True,
    algorithm='lloyd')
```

To order topic clusters, we used the 'fractional ordering' of the sentences, namely, let d be a document with length n , then the fractional ordering f of sentence i at position i in d is:

$$f(i, d) = \frac{i}{n}$$

For example, the first sentence in the given document is always $\frac{1}{n}$, the second $\frac{2}{n}$, etc. We divide by the number of sentences in the document to try to normalize both short and long documents.

We then ordered each topic cluster t_k with m sentences by their respective median fractional ordering, namely, for each sentence j in the given

topic, the median fractional ordering f_{med} is found by sorting all the fractional orderings for each sentence, and using the one in the middle.

Using average (see below) versus median fractional ordering was another hyperparameter we chose from. Please see Table 1 for an error analysis between them.

The average fractional ordering f_{avg} is:

$$f_{avg}(t_k) = \frac{\sum_j f(j, d)}{m}$$

We then order the sentences in the given summary based on which topic cluster they appear in and the fractional ordering of that topic. In other words, if $f_{med}(t_k) < f_{med}(t'_k)$, then any sentence that appeared in the summary and in topic cluster t_k would appear before all the sentences that appeared in the summary and in topic cluster t'_k .

It is possible that multiple sentences appear in the given summary and the same topic cluster. In this case, we order the sentences by their own fractional ordering. In other words, for sentence i and sentence j in documents d and d' respectively, if $f(i, d) < f(j, d')$ we put sentence i before sentence j in the ordered summary.

4.3.2 Zero-shot Learning

We trained a zero-shot learning language model based on the implementation of Reorder-BART (RE-BART) by Chowdhury et al.. RE-BART is a fine-tuned model based on BART by to identify a coherent order for a given set of shuffled sentences. We shuffle the sentences of each input document set based on the Gap sentence generation content selection method without masking. We mark the index of the sentences at the beginning of each sentences. The model takes the sets of shuffled sentences with sentence-specific markers as input and generates a sequence of position markers of the sentences in the ordered text. We trained the model using Huggingface transformer library Wolf et al. (2020). We trained the model using the PyTorch framework with a NVIDIA A100 GPU. We trained the model with batch size of 4 and epoch of 24.

4.3.3 ROUGE score ranking

Due to the input size limitation for the majority of the language model, we have to truncate the input text to 1024 tokens. After we mask the important sentence, we then use the ROUGE score ranking calculated for the gap sentence generation to discard sentences that are ranked in the low thirty

percent. We keep the ordering of the remainder of the sentences. Discarding unimportant sentences based on ROUGE score helps including more important sentences from multiple documents. When discarding, we calculate the token length for each added sentences and stop when adding additional sentence will cause the token size to exceed 1024 tokens. This make sure we have full sentences for the input sequence.

To improve the system, we experiment on multiple parameters, such as what the percentage of the sentences should we discard. We also calculate the ROUGE score ranking based on either the average score based on the gold summary or based on all of the remaining sentences in all provided documents.

4.4 Content Realization

4.4.1 Entity-driven Rewrite

We developed an entity replacement method inspired by the noun phrase rewriting method for multi-document summarization method described in Siddharthan et al. (2011). We take each document set being analyzed and cluster spans into categories denoting the same entity. We accomplish this task using an experimental module spaCy (Honnibal et al., 2020). The module obtains token embeddings from RoBERTa and then obtains a score for embedding pairs before passing the score to a linear classifier (along with other features) to determine whether the spans refer to the same entity (Kádár et al., 2022). Because the number of spans is $O(n^2)$, calculating each potential coreference is computationally expensive, so the spaCy implementation prunes the number of spans to compare before obtain the results from the linear classifier, adapting an algorithm described in (Dobrovolskii, 2021).

We adapted our NP replacement algorithm from Siddharthan et al. (2011). We use our ILP and LexRank methods to extract the sentences with the highest weights. We also obtain coreference clusters for the entire document set (concatenating the documents) using spaCy’s module described above. We apply the model to the concatenated document to minimize the number of clusters with the same or similar spans. Then, using, spaCy base module (Honnibal et al., 2020), we obtain the noun phrases contained in the highest-ranked sentence $\{NP | NP \in s_0\}$. For each NP in the highest-ranked sentence, we examine its corefer-

Table 1: Error Analysis for summary ordering using Topic Clustering

docset	Unordered	Ordered
D1001-A	The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in. So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell. Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart. But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.	So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell. Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart. The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in. But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.
D1002-A	Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade. Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic. While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases. They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.	Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade. Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic. While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases. They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.

ring NPs in the document set and it with the longest pre-modifying NP. We use a rough heuristic to extract the longest pre-modified by discard all text followed by appositives, since those formed the bulk of the post-modifiers described in Siddharthan et al. (2011). If the longest pre-modifying NP is the same length as the NP itself, we look for the longest post-modifying phrase. The cluster containing the longest pre-modifying NP is indexed. For all future NPs and sentences, if $NP_i \in s_j$ is contained in a seen cluster c_k , then the shortest non-pronominal NP is chosen to replace it. If the $NP_i \in s_j$ is not contained in a seen cluster c_k , it is replaced with the longest modifying phrase. This intuition was drawn from the observations of Siddharthan et al. (2011), about keeping NPs of previously seen referents terse and expanding upon unseen referents.

In addition, added a few small rules to ensure a readable summary. We did not replace the pronoun "I" in the context of quotes. We also did not replace

referents with other pronouns to avoid issues with case (e.g., "we" vs. "our").

4.4.2 Redundancy Removal

After obtaining a ranked order of extracted sentences, we experimented with removing redundant sentences using the Jaccard similarity measure. Sentences with the top- k values in ranking are extracted and added to the summary text until the summary reaches the maximum length. To avoid repeated information, we use the Jaccard similarity measure (provided by NLTK) to calculate each sentence's similarity to sentences already included in the bibliography (Bird and Loper, 2004). To calculate Jaccard similarity for sentences, let J denote the similarity function, w denote a word and s_i, s_j denote two sentences.

$$J(s_i, s_j) = \frac{|\{w|w \in s_j\} \cap \{w|w \in s_i\}|}{|\{w|w \in s_i\} \cup \{w|w \in s_j\}|} \quad (17)$$

Additionally, sentences that are too long are discarded. We discard long sentences because we

believe these sentences are too likely to be too information-dense to be useful for constructing a summary. We also discard sentences that were a part of a quotation that was cut off by the tokenizer since we currently have no way of evaluating the importance of the sentence combine with its reconstructed context. The Redundancy removal method was employed by the LexRank algorithm, but not by others.

4.4.3 Large Language Model Training

We trained our model based on "google/pegasus-large" and "google/pegasus-cnn_dailymail". We utilized the Huggingface's transformers library for the experiment. The experiments are conducted in PyTorch framework using NVIDIA Tesla A100 GPU. We trained on the training data and use the devtest data to select the best model based on the ROUGE 1 score of output summary based on devtest data.

5 Results

Table 2: ROUGE Recall Scores based on **devtest** data files

	ROUGE1	ROUGE2
Best of TAC 2011		0.09574
Binary ILP	0.33320	0.07401
LexRank	0.21925	0.05966
GSG LLM	0.31355	0.08191
Baseline Top K	0.26193	0.05846

Table 3: ROUGE Recall Scores based on **evaltest** data files

	ROUGE1	ROUGE2
Best of TAC 2011		0.13440
Binary ILP	-	-
LexRank	0.21925	0.05966
GSG LLM	0.29880	0.06799
Baseline Top K	-	-

5.1 Large Language Model

For our experiment for finetuning the PEGASUS model, we tested difference combination of the training arguments, such as different epoch and difference batch sizes. Table 4 listed all parameter changes we make for out experiment. ROUGE-on stands for calculating the ROUGE score of a selected sentence from a document, either with the other sentences in that document (single) or with

the other sentences in all the documents in a docset (multi). We also experiment on different epoch. Discard stands for the percentage of sentences we selected to discard when generating the input. For example, if we change the ROUGE-on parameters to multi, then choosing 50% discard rate means that we will discard the bottom 50% of the sentences in all of the documents in a docset based on how high the ROUGE score is. Combine masking stands for whether or not we concatenate the gap-sentences into a single mask token. The result shows that when we calculate the ROUGE score of the selected sentence with the rest of the sentences in all documents in one docset, discard bottom 50% of the sentences, and combine multiple masked sentences into a single mask token achieve the highest ROUGE 1 and ROUGE 2 score for the devtest data. We also experiments with both "google/pegasus-large" and "google/pegasus-cnn_dailymail" checkpoints for the PEGASUS model. We find out that the news-summarization focused "google/pegasus-cnn_dailymail" performs better than "google/pegasus-large", which trained on more datasets than "google/pegasus-cnn_dailymail".

We also experiment on the zero-shot learning method for information ordering using the BART model. The result, however, is disappointing. The model we trained for our experiment failed to generate a sequence of position makers as output. Instead, the model directly generates the ordered text as the output where each sentences are compressed and no longer contain the same level of information compare to the original sentences. We tried to investigate the cause of the issue by reducing input sizes, and increase training data coverage by including all of the sentences from training, testing, and validation data into groups of six sentences. We tried to tweak other parameters such as increasing the training epoch from 8 to 24. However, none of the above methods leads to having the generated output from the fine-tuned model produces sentence indexes, and we choose not to move forward with this information ordering method.

5.2 Ablation on Top Model Implemented: ILP

We chose the ILP model with hyperparameters as follows as the top model

Min Sent Length = 25
n-gram = Unigram
delta tf = 0.01

model	ROUGE-on	epoch	Discard	Combine Masking	ROUGE1	ROUGE2
pegasus-large	single	6	50%	True	0.21037	0.06214
pegasus-large	multi	12	50%	True	0.26419	0.05367
pegasus-large	multi	24	50%	True	0.28415	0.06464
pegasus-cnn_dailymail	multi	24	50%	True	0.31355	0.08191
pegasus-large	multi	12	30%	True	0.24330	0.04773
pegasus-large	multi	12	30%	False	0.24263	0.05343

Table 4: Large language model ROUGE Recall Scores with different parameters

Exp-ID	Min Sent Length	n-gram	delta tf	delta idf	Elim Punc	Lower-casing	log	ROUGE1	ROUGE2
J0	25	Unigram	0.01	0.7	No	Yes	Yes	33.32	7.41
J1	None	Unigram	0.01	0.7	No	Yes	Yes	32.508	6.849
J2	25	Bigram	0.01	0.7	No	Yes	Yes	30.059	7.016
J3	25	Trigram	0.01	0.7	No	Yes	Yes	27.682	6.078
J4	26	Unigram	0.001	0.7	No	Yes	Yes	33.114	7.349
J5	25	Unigram	0.01	0.001	No	Yes	Yes	22.592	3.107
J6	25	Unigram	0.01	0.7	Yes	Yes	Yes	33.32	7.405
J7	25	Unigram	0.01	0.7	No	No	Yes	33.32	7.405
J8	25	Unigram	0.01	0.7	No	Yes	No	33.32	7.415

Table 5: The results of the experiments that we ran for our ablation test of our top-model ILP. The top row is the top-model with the best combination of hyperparameters that gets us the our very best ROUGE1 score.

delta idf = 0.7
 Eliminate Punctuation = No
 Lower Casing = Yes
 log = Yes

We performed an ablation test on this top model, to see which hyperparameter causes the greatest increase in ROUGE score. As we see in table 5, decreasing the *delta idf* for smoothing to around 0, causes the greatest decrease in ROUGE score, (-10.728). It also appears that our choice of n-gram has an impact on the system with unigrams performing the best (-0.0), and trigrams performing the worst (-5.638). Interestingly, we found that discarding sentences under a certain sentence length has very little effect on the performance (-0.812). Likewise, whether we choose to log the *tf-idf* values, eliminate all punctuation, choose to lowercase all terms, or choose a *delta tf* that is close to 0, these hyperparameters had little to no effect on the performance of the system.

6 Discussion

6.1 Error Analysis of *devtest* data

We perform a casual error analysis for the summaries based on **devtest** docset D1006, which are shown in Table 6.

The improved Binary ILP method makes leaps in achieving a better summary from D3. The old summary seemed to give important facts concerning the FDA and withdrawing Vioxx, but seemed to have a lot of more "unimportant facts" that didn't help the reader get a clear idea of what the article is about. An example of this is mentioning about rewriting abstract conclusions, and the amount of teleconferences that were gone to. The improved summary gives a clear picture of the latest update on Vioxx and the new findings of effects its gives. Comparing the output of D4 and D5, there does not seem to be a difference. This may be because the new hyperparameters from clustering did not effect the sentence ordering here, and content realization did not find anything to use co-reference resolution on.

The improved LexRank method no longer produce unnecessary information such as website ad-

dress. The improved method successfully mention Vioxx, where the old method did not. However, the improved method still failed to catch one of the core story point, that Vioxx is recalled by the company. The improved method does capture that Vioxx has potential cardiovascular risks. Interestingly, the using TF-IDF vectors resulted in better performance than using word2vec or DistilBERT sentence vectors. A comparison is shown in Table 10 in Appendix B.

The summary produced by the improved GSG LLM method provides more specific details than the previous system iteration. For example, it mentions that Vioxx was used by 20 million Americans, was Merck’s top-selling product, and that Merck had spent \$195 million to promote it. It also mentions that the FDA had been concerned about the drug’s cardiovascular risks since at least 2000 but did not issue a warning until 2004. These details provide more context and a better understanding of the situation. Another improvement is that it provides a clearer timeline of events. It mentions that Vioxx was approved in 1999, that the FDA had been concerned about its cardiovascular risks since at least 2000, and that the drug was recalled in 2004. This helps the reader understand the sequence of events and the time frame in which they occurred. Compared to the gold summary, however, the improved system still lacks certian information. For example, the improved GSG LLM method generated summary lacks information about the specific clinical trial that led to the recall of Vioxx. The gold summary mentions that the clinical trial was for the use of Vioxx in colon cancer and that it showed unacceptable rates of stroke and heart attack. The generated summary also does not mention that Vioxx was a COX inhibitor, which was safer for the digestive tracts of arthritis patients. Additionally, the gold summary mentions concerns about drug manufacturers’ advertising and the FDA’s role in ensuring the safety of drugs on the market, which is not mentioned in the generated summary.

The summary produced by the baseline mentions a small amount of important details such withdraw the specific drug and how many people used it. But fails to mention many other important details mentioned in the gold standard and by the other methods.

6.2 Error Analysis of *evaltest* data

We perform a casual error analysis for the summaries based on **evaltest** docset D1105, which are shown in Table 7.

For the ILP method, it mentions most of the details in the gold standard summary. It does miss the exact date of the crash, and how many killed/possible survivors. It also missed the detail on the unknown cause of the crash, but does mention about the stormy weather.

For the GSG LLM method, the generated summary correctly states that an Adam Air Boeing 737-400 plane with 102 people on board crashed in a mountainous area near the town of Polewali, on its way from Surabaya to Manado. It also correctly states the casualty count of the crash. These are some of the main points of the gold text that the generated summary accurately captures. However, the generated summary lacks additional details such as the weather condition might be a factor of the crash, and there were three Americans on board. Lastly, the generated summary introduces new information about the Indonesian Navy sending planes to carry the bodies of its members, which is not mentioned in the gold text. Overall, the generated summary captures some of the main points of the gold text, but also contains errors and omissions.

The baseline seems to mention mostly about a different event, but mentions the main event the gold standard summary is concerned about in the last sentence. This may be because the first article in the docset mentions about both the events in the first few lines.

6.3 Error analysis of Content Realization

Will appear in final submission.

7 Conclusion

In conclusion, our paper explores multiple approaches to multi-document summarization, including both extractive and abstractive methods. We have built end-to-end systems using the selected methods and benchmarked them to evaluate their effectiveness. Our error analysis provides insights into the strengths and weaknesses of the selected methods, paving the way for future research in this field.

8 Appendices

A Workload distribution

A.1 D1 Workload

- Anna Batra set up the Github repository, turned in D1
- Junyin Chen got the team together and set up a communication channel
- Sam Briggs set up the Overleaf file and sent out a when-to-meet to schedule weekly meetings
- Hilly Steinmetz edited the Overleaf file to prepare it for D1.

A.2 D2 Workload

- Anna Batra and Sam Briggs wrote test code to test the file structure of the output docSets, created the outline for the presentation, and updated the report.
- Junyin Chen wrote code for tokenizing documents in the docSets using spaCy, PR reviewed the code to merge with Hilly's, cleaned up the code, and created slides for the pre-processing section.
- Hilly Steinmetz wrote the code for the pre-processing steps before tokenization, such as locating paths for AQUAINT and AQUAINT2 files. Hilly also wrote code for tokenizing documents using NLTK.

A.3 D3 Workload

- Anna Batra and Sam Briggs wrote the code to create a json file to easily access our data for the rest of the project. They also wrote the code for the TF-IDF and Linear Programming content selection methods. The Linear Programming information ordering and content realization was also written by them. They also drew the system architecture.
- Junyin Chen wrote the code to create JSON file writer which contains doc_id, text for summarization, the gold standard summarization based on doc_id for both docsetA and docsetB. The writer help cache the JSON file for easier access. He also wrote the code for Gap sentences generation content selection method, truncate the input text based on ROUGE score

for information ordering, and write the code for training a large language model for content realization. He also performs quick error analysis.

- Hilly Steinmetz wrote the code for the LexRank method and debugged issues with the original XML document parser.

Everyone worked on the paper and presentation slides for the parts we explicitly worked on.

A.4 D4 Workload

- Anna Batra and Sam Briggs incorporated ngrams into TF-IDF and helped test the IDF to work over the entire data. They also worked on improving ILP and wrote the clustering information ordering method. They updated the system architecture.
- Junyin Chen wrote and trained zero shot information ordering language model. He also improved pre-processing scripts for data parsing, such as using different ROUGE score and comparing with different documents. He done multiple experiments with information ordering model and summarizing model with multiple parameter combination.
- Hilly Steinmetz created an interface to generate vectors using TF-IDF, Word2vec, or DistilBERT. He also worked on creating a new TF-IDF class that takes various inputs to modify its behavior (e.g., smoothing). Lastly, he implemented the improvements to LexRank.

Everyone worked on the paper and presentation slides for the parts we explicitly worked on.

A.5 D5 Workload

- Anna Batra tested the clustering information ordering method on different hyperparameters using an error analysis and ran ablation on the top model. She also wrote the baseline top K and redrew the system architecture. She updated this all into the slides.
- Sam Briggs fixed the paper based on the feedback by Fei, incorporated what Anna did into the paper, and overall worked to improve the quality of the paper.
- Junyin Chen experimented on zero-shot information ordering methods. He wrote the abstract, introduction, and related work section

of the paper. He also updated the correspondent section in the paper about LLM.

- Hilly Steinmetz implemented the entity-driven rewrite and experimented with LexRank to improve its results (with marginal improvements).

Table 10: LexRank Experiments (ROUGE F-Scores)

Vector	R-1	R-2
TF-IDF	0.231	0.060
Word2Vec	0.175	0.036
DistilBERT	0.178	0.042

B Code repository and additional software and data used in your system

B.1 Source Code

The repository for our project can be found on Github at github.com/LING-575-Summarization/Summarization.

B.2 Packages

We used the following packages for our system:

B.2.1 Pre-processing

- lxml.etree (for processing AQUAINT, AQUAINT2, TAC files)
- xml.etree.ElementTree (for processing doc-SetA file lists)
- spaCy 2.0 (for word tokenization on paragraphs)
 - English model “en_core_web_sm”
- NLTK
 - English model ‘tokenizers/punkt/english.pickle’ (for sentence and word tokenization)
 - nltk.util.ngrams (for TF-IDF ngrams)

B.2.2 Content Selection

- PuLP (for binary ILP)
- rouge-score
- Datasets
- Evaluate
- Pytorch
- Transformers
- Anaconda (for virtual environment)

B.2.3 Information Ordering

- sci-kit learn (for Topic Clustering)

C Figures and Tables

C.1 Ablation Test

C.2 LexRank Experiments

C.3 ROUGE score from previous Deliveries

Table 11: ROUGE Recall Scores

	ROUGE1	ROUGE2
Binary ILP (D4)	0.33697	0.07437
Binary ILP (D3)	0.12085	0.01533
LexRank (D4)	0.21925	0.05966
LexRank (D3)	0.13720	0.02341
GSG LLM (D4)	0.26419	0.05367
GSG LLM (D3)	0.21037	0.06214

References

- R. "Barzilay, N. Elhadad, and K." McKeown. 2002. [Inferring strategies for sentence ordering in multi-document news summarization](#). *Journal of Artificial Intelligence Research*, 17(17):35–55.
- Regina Barzilay, Kathleen R. McKeown, and Michael Elhadad. 1999. [Information fusion in the context of multi-document summarization](#). In *Proceedings of the 37th Annual Meeting of the Association for Computational Linguistics*, pages 550–557, College Park, Maryland, USA. Association for Computational Linguistics.
- Steven Bird and Edward Loper. 2004. [NLTK: The natural language toolkit](#). In *Proceedings of the ACL Interactive Poster and Demonstration Sessions*, pages 214–217, Barcelona, Spain. Association for Computational Linguistics.
- Jaime Carbonell and Jade Goldstein. 1998. [The use of mmr, diversity-based reranking for reordering documents and producing summaries](#). In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, SIGIR '98, page 335–336, New York, NY, USA. Association for Computing Machinery.
- Asli Celikyilmaz, Antoine Bosselut, Xiaodong He, and Yejin Choi. 2018. [Deep communicating agents for abstractive summarization](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1662–1675, New Orleans, Louisiana. Association for Computational Linguistics.
- Jianpeng Cheng and Mirella Lapata. 2016. [Neural summarization by extracting sentences and words](#). In

949	<i>Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 484–494, Berlin, Germany. Association for Computational Linguistics.	
950		
951		
952		
953	Sumit Chopra, Michael Auli, and Alexander M. Rush. 2016. Abstractive sentence summarization with attentive recurrent neural networks . In <i>Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 93–98, San Diego, California. Association for Computational Linguistics.	
954		
955		
956		
957		
958		
959		
960		
961	Somnath Basu Roy Chowdhury, Faeze Brahman, and Snigdha Chaturvedi. 2021. Is everything in order? a simple way to order sentences .	
962		
963		
964	Arman Cohan, Franck Dernoncourt, Doo Soon Kim, Trung Bui, Seokhwan Kim, Walter Chang, and Nazli Goharian. 2018. A discourse-aware attention model for abstractive summarization of long documents . In <i>Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)</i> , pages 615–621, New Orleans, Louisiana. Association for Computational Linguistics.	
965		
966		
967		
968		
969		
970		
971		
972		
973		
974	Vladimir Dobrovolskii. 2021. Word-level coreference resolution . In <i>Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing</i> , pages 7670–7675, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.	
975		
976		
977		
978		
979		
980	Günes Erkan and Dragomir R Radev. 2004. Lexrank: Graph-based lexical centrality as salience in text summarization. <i>Journal of artificial intelligence research</i> , 22:457–479.	
981		
982		
983		
984	Alexander Fabbri, Irene Li, Tianwei She, Suyi Li, and Dragomir Radev. 2019. Multi-news: A large-scale multi-document summarization dataset and abstractive hierarchical model . In <i>Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics</i> , pages 1074–1084, Florence, Italy. Association for Computational Linguistics.	
985		
986		
987		
988		
989		
990		
991	Kavita Ganesan, ChengXiang Zhai, and Jiawei Han. 2010. Opinosis: A graph based approach to abstractive summarization of highly redundant opinions . In <i>Proceedings of the 23rd International Conference on Computational Linguistics (Coling 2010)</i> , pages 340–348, Beijing, China. Coling 2010 Organizing Committee.	
992		
993		
994		
995		
996		
997		
998	Sebastian Gehrmann, Yuntian Deng, and Alexander Rush. 2018. Bottom-up abstractive summarization . In <i>Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing</i> , pages 4098–4109, Brussels, Belgium. Association for Computational Linguistics.	
999		
1000		
1001		
1002		
1003		
	Dan Gillick, Benoit Favre, and Dilek Hakkani-Tur. 2008. The icsi summarization system at tac 2008. <i>Theory and Applications of Categories</i> .	1004 1005 1006
	Aria Haghighi and Lucy Vanderwende. 2009. Exploring content models for multi-document summarization . In <i>Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics</i> , pages 362–370, Boulder, Colorado. Association for Computational Linguistics.	1007 1008 1009 1010 1011 1012 1013
	Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. 2020. spaCy: Industrial-strength Natural Language Processing in Python .	1014 1015 1016
	Dan Jurafsky and James H. Martin. 2022. Speech and Language Processing , 3rd edition (draft) edition. Online.	1017 1018 1019
	Ákos Kádár, Paul O’Leary McCann, Edward Schmuhl, Sofie Van Landeghem, Adriane Boyd, Madeesh Kannan, and Victoria Slocum. 2022. End-to-end Neural Coreference Resolution in spaCy .	1020 1021 1022 1023
	Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. 2019. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension .	1024 1025 1026 1027 1028
	Wencan Luo, Fei Liu, Zitao Liu, and Diane Litman. 2018. A novel ilp framework for summarizing content with high lexical variety . <i>Natural Language Engineering</i> , 24(6):887–920.	1029 1030 1031 1032
	Kathleen McKeown and Dragomir R Radev. 1995. Generating summaries of multiple news articles. In <i>Proceedings of the 18th annual international ACM SIGIR conference on Research and development in information retrieval</i> , pages 74–82.	1033 1034 1035 1036 1037
	Rada Mihalcea and Paul Tarau. 2004. TextRank: Bringing order into text . In <i>Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing</i> , pages 404–411, Barcelona, Spain. Association for Computational Linguistics.	1038 1039 1040 1041 1042
	Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. Efficient estimation of word representations in vector space. <i>arXiv preprint arXiv:1301.3781</i> .	1043 1044 1045 1046
	Ramesh Nallapati, Bowen Zhou, Cicero dos Santos, Caglar Gulcehre, and Bing Xiang. 2016a. Abstractive text summarization using sequence-to-sequence RNNs and beyond . In <i>Proceedings of the 20th SIGNLL Conference on Computational Natural Language Learning</i> , pages 280–290, Berlin, Germany. Association for Computational Linguistics.	1047 1048 1049 1050 1051 1052 1053
	Ramesh Nallapati, Bowen Zhou, and Mingbo Ma. 2016b. Classify or select: Neural architectures for extractive document summarization .	1054 1055 1056

- Shashi Narayan, Shay B. Cohen, and Mirella Lapata. 2018. [Ranking sentences for extractive summarization with reinforcement learning](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1747–1759, New Orleans, Louisiana. Association for Computational Linguistics.
- Karolina Owczarzak and Hoa Trang Dang. 2009. Overview of the tac 2009 summarization track. In *Proceedings of the Text Analysis Conference (TAC)*.
- Karolina Owczarzak and Hoa Trang Dang. 2010. Overview of the tac 2010 summarization track. In *Proceedings of the Text Analysis Conference (TAC)*.
- Karolina Owczarzak and Hoa Trang Dang. 2011. Overview of the tac 2011 summarization track: Guided task and aesop task. In *Proceedings of the Text Analysis Conference (TAC 2011)*, Gaithersburg, Maryland, USA, November.
- Lawrence Page, Sergey Brin, Rajeev Motwani, and Terry Winograd. 1999. The pagerank citation ranking: Bringing order to the web.
- Romain Paulus, Caiming Xiong, and Richard Socher. 2017. [A deep reinforced model for abstractive summarization](#).
- Dragomir R. Radev, Hongyan Jing, and Malgorzata Budzikowska. 2000. [Centroid-based summarization of multiple documents: sentence extraction, utility-based evaluation, and user studies](#). In *NAACL-ANLP 2000 Workshop: Automatic Summarization*.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2019. [Exploring the limits of transfer learning with a unified text-to-text transformer](#).
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *ArXiv*, abs/1910.01108.
- Abigail See, Peter J. Liu, and Christopher D. Manning. 2017. [Get to the point: Summarization with pointer-generator networks](#).
- Yohei Seki. 2003. [Sentence extraction by tf/idf and position weighting from newspaper articles](#). *National Institute of Informatics: Proceedings of the Third NTCIR Workshop*.
- Advaith Siddharthan, Ani Nenkova, and Kathleen McKeown. 2011. [Information status distinctions and referring expressions: An empirical study of references to people in news summaries](#). *Computational Linguistics*, 37(4):811–842.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander Rush. 2020. [Trans-formers: State-of-the-art natural language processing](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online. Association for Computational Linguistics.
- Jingqing Zhang, Yao Zhao, Mohammad Saleh, and Peter J. Liu. 2019. [PEGASUS: pre-training with extracted gap-sentences for abstractive summarization](#). *CoRR*, abs/1912.08777.

Table 6: Generated summaries based on **devtest** docset D1006

gold	On Sept 30, Merck voluntarily recalled the pain killer Vioxx, used by almost 2 million, after clinical trials for its use in colon cancer showed unacceptable rates of stroke/heart attack. Results corroborated earlier warnings that had not resulted in recalls by the Food and Drug Administration (FDA). As a COX inhibitor, Vioxx was safer for digestive tracts, important for arthritis patients. Merck's advertising campaigns did not clearly warn about side effects. The case highlighted concerns about drug manufacturers' advertising and FDA's role in insuring safety of drugs on the market. Safety of other COX inhibitors is now a concern."
ILP D5	Merck officials said last week its latest research showed an increased risk of heart attack and other cardiovascular complications in patients who took Vioxx for at least 18 months. Heavily advertised as an arthritis drug, Vioxx was pulled from the market last week after its maker said a study showed it doubled the risk of heart attack and stroke. But some doctors say this group of drugs may work in a way that increases the risk of heart problems for some patients, and they point to this latest information as additional reason for concern.
ILP D4	Merck officials said last week its latest research showed an increased risk of heart attack and other cardiovascular complications in patients who took Vioxx for at least 18 months. Heavily advertised as an arthritis drug, Vioxx was pulled from the market last week after its maker said a study showed it doubled the risk of heart attack and stroke. But some doctors say this group of drugs may work in a way that increases the risk of heart problems for some patients, and they point to this latest information as additional reason for concern.
LexRank D4	With Vioxx, researchers had been warning about the drug's possible cardiovascular risks since 2000, only a year after it was approved by the FDA . Data from a company study found then that users had four times as many heart attacks and strokes as those who used another painkiller . But the data was not definitive, and Merck, which even critics say is one of the most responsible drug companies, repeatedly reassured the medical and financial communities that Vioxx was safe.
LexRank D3	"That's the tragedy here . " And if courts determine that Merck was negligent, the company will pay a heavy price in compensation. It is in the insurance industry's interest, the FDA's interest and the federal government's interest—because the federal government is a major provider of health insurance—either to require drug companies to conduct such comparative tests or to set up a neutral agency to do so. FDA: http://www.fda.gov/
GSG LLM D5	In September 2004, Merck & Co. recalled its arthritis drug Vioxx after a clinical trial showed it doubled the risk of heart attacks and strokes. The drug had been used by 20 million Americans since its approval in 1999 and was the company's top-selling product. Merck had spent \$195 million to promote Vioxx as a wonder drug for the aging baby boomers. The FDA, which approved Vioxx for use, had been concerned about the drug's cardiovascular risks since at least 2000 but did not issue a warning until 2004.
GSG LLM D4	Merck recalled Vioxx in September 2004 after a study showed that it doubled the risk of heart attacks and strokes in older people taking it for at least three years. The drug had been approved by the FDA in 1999 for arthritis. Merck had promoted Vioxx as a way to lower blood pressure and cholesterol, but the study showed that it increased the risk of heart attacks and strokes. Merck's decision to withdraw Vioxx from the market raised questions about aggressive marketing of the drug before its long-term safety had been proven.
Baseline Top K (D5)	Merck Co. said its surprise decision Thursday to withdraw the arthritis drug Vioxx—used by about 2 million people worldwide—was driven by recent evidence that the drug's adverse side effects outweighed any potential benefits. But that wasn't really news.

Table 7: Generated summaries based on **evaltest** docset D1105

gold	Boeing 737-400 plane with 102 people on board crashed into a mountain in the West Sulawesi province of Indonesia, on Monday, January 01, 2007, killing at least 90 passengers, with 12 possible survivors. The plane was Adam Air flight KI-574, departing at 12:59 pm from Surabaya on Java bound for Manado in northeast Sulawesi. The plane crashed in a mountainous region in Polewali, west Sulawesi province. There were three Americans on board, it is not know if they survived. The cause of the crash is not known at this time but it is possible bad weather was a factor.
ILP	An Indonesian lawmaker Tuesday criticised what he said was the slow deployment of search and rescue teams to a mountainous area of Sulawesi island where a plane carrying 102 people crashed. An Indonesian passenger plane carrying 102 people disappeared in stormy weather on Monday, and rescue teams were sent to search an area where military aviation officials feared the Boeing 737-400 aircraft may have crashed. An Adam Air Boeing 737-400 plane with 102 people on board crashed in a mountainous area near the town of Polewali late Monday on its way from Surabaya to Manado.
LexRank	
GSG LLM	The Indonesian Navy (TNI AL) has sent two Cassa planes to carry the bodies of five of its members who were killed in a plane crash in Sulawesi late Monday. An Adam Air Boeing 737-400 plane with 102 people on board crashed in a mountainous area near the town of Polewali late Monday on its way from Surabaya to Manado. At least 90 people, including five TNI AL members, were killed in the crash.
Baseline	Indonesian President Susilo Bambang Yudhoyono said Tuesday he was deeply concerned with the crash of a passenger plane and the sinking of a ferry in the last few days that might have killed hundreds of people. Earlier on Friday, a ferry carrying 628 people sank off the Java coast with some 400 passengers reported missing. A Boeing 737-400 plane with 102 people onboard crashed into a mountain in the West Sulawesi province Monday, killing at least 90 people. JAKARTA, Jan. 2 (Xinhua)

Docset	TF-IDF	Word2Vec	DistilBert
D1001-A	<p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p> <p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p>	<p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p>	<p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p> <p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p>
D1002-A	<p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p> <p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p>	<p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p> <p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p>	<p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p> <p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p>

Table 8: Error Analysis for summary ordering using Topic Clustering using **mean** Fractional Ordering

Docset	TF-IDF	Word2Vec	DistilBert
D1001-A	<p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p>	<p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p>	<p>Graham praised the Columbine community for uniting under the pain of a tragedy that could have torn it apart.</p> <p>But Wells said he is more interested in simply trying to have fun and move beyond the tragedy that put his life on hold.</p> <p>So many forms of community, rippling outward from Columbine High and across the planet, have come together since last week's violence that it was difficult to tell.</p> <p>The school wanted to make sure there was enough to eat since students couldn't leave campus for lunch and get back in.</p>
D1002-A	<p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p> <p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p>	<p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p> <p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p>	<p>They are accused of firing 41 times at Amadou Diallo while searching for a rape suspect on Feb. 4.</p> <p>Police officers in criminal trials have often asked for a judge to decide their case, fearing that juries would be unsympathetic.</p> <p>While the trial date would come nearly a year after Diallo's death on the night of Feb. 4, it is not unusual in such high-publicity cases.</p> <p>Several of the officers are said to have told associates that they continued firing because Diallo did not fall even after they had unleashed the fusillade.</p>

Table 9: Error Analysis for summary ordering using Topic Clustering using **median** Fractional Ordering